

# N Shashi Kumar

Hyderabad | Ph: +91 7989265399 | Email: [Shashinarmala29@gmail.com](mailto:Shashinarmala29@gmail.com) | [Github](#) | [Linkedin](#)

## OBJECTIVE

Motivated and detail-oriented fresher aspiring to start a career as a DevOps Engineer. Eager to apply foundational knowledge of CI/CD, cloud platforms, Linux, and scripting to support efficient software delivery. Passionate about learning modern DevOps tools and practices to help teams build, deploy, and scale applications reliably

## EDUCATION

<b>Avanthi Institute of Engineering and Technology</b>	<b>Hyderabad</b>
<b>Bachelors in Computer Science Engineering – Data Science</b>	<b>SEP 2021 – 2025</b>
<ul style="list-style-type: none"><li><b>Relevant Coursework:</b> Software Engineering , Machine Learning , Computer Networks , Database Management Systems , Cloud Computing, Data Mining .</li></ul>	

## EXPERIENCE

DevOps Intern	Sep 2025 – Dec 2025
<b>Trav</b>	<b>Remote, India</b>
<ul style="list-style-type: none"><li>Created reusable GitHub Actions workflow templates for Python microservices, enabling faster CI/CD setup, standardization, and easy inheritance across multiple services.</li><li>Researched and estimated cloud costs across Azure and AWS, comparing service pricing and scaling strategies to design a cost-efficient and scalable cloud architecture for the frontend application.</li><li>Defined ECS task definitions and configured AWS networking components, IAM roles, and supporting cloud services to deploy the application in a secure production-ready environment.</li><li>Developed Custom cloud watch dashboards implementing <b>2x predictive</b> monitoring and observability of AWS ECS services and Troubleshooted production errors</li><li>Managed and optimized infrastructure for a production application receiving <b>11k–12k daily requests</b>, ensuring stable performance and zero downtime during peak hours.</li><li>Implemented strict branch protection and strategy rules like trunk based development and made concrete SDLC for application CICD</li></ul>	

Junior DevOps Engineer – Contract	Feb 2025 – July 2025
<b>Vtex.AI</b>	<b>Remote, India</b>
<ul style="list-style-type: none"><li>Developed robust automated testing pipelines using GitHub Actions, integrating SAST (Static Application Security Testing), UAT (User Acceptance Testing), and regression testing to maintain high code quality throughout the CI/CD process.</li><li>Implemented Canary deployments with automated rollback mechanisms, ensuring zero-downtime releases and enabling safe, progressive feature rollouts using version tagging.</li><li>Streamlined automated GitHub releases and semantic versioning across a multi-repository</li></ul>	

architecture, improving deployment traceability and collaboration.

- Designed and implemented a production-ready CI/CD pipeline using Jenkins to automate the build and deployment of a Python-based application on AWS ECS, reducing manual effort by over 90%.
- Built a fully containerized architecture with Docker and Amazon ECR, optimizing image size and ensuring consistent deployment across environments.
- Provisioned cloud infrastructure using Terraform, automating tasks such as ECS cluster setup, IAM role creation, networking, and CloudWatch log configuration—following Infrastructure as Code (IaC) best practices.
- Integrated CloudWatch Logs and ECS service monitoring, enhancing system observability and reducing issue resolution time by 70%.
- Achieved an 80% reduction in deployment time while maintaining zero downtime, significantly improving the release process and user experience.

### Kitikiplot

March 2025

- Kitikiplot is a python open source library with more than 6K+ Installations and Recognized by FOSS
- Developed Python Unit test cases using Pytest Ensuring Code Quality
- Implemented a Automated Github Actions Pipelines for Automated Tests

---

## PROJECTS

---

### **Finch**

SEP 2025

- Built an end-to-end CI/CD pipeline from source code to production deployment, implementing Docker multi-stage builds, image registry workflows, unit testing, Kubernetes deployments, and automatic scaling using HPA.
- Established a secure RDS connection using Cloud SQL Auth Proxy, enabling encrypted communication between GKE workloads and the database.
- Implemented observability using Prometheus and Grafana with custom alerts to track cluster health, performance, and application behavior.
- Successfully predicted potential production failures and proactively resolved issues using monitoring insights, improving system stability and reliability.

### **BeanClassifier KF**

NOV 2025

- Built an end-to-end ML training and serving workflow using the **Beans dataset** (three-class leaf disease classification) and automated the entire process with **Kubeflow Pipelines**.

- Designed modular Kubeflow components (data preprocessing , training ,evaluation , model exporting) and containerized each stage using Docker for fully reproducible pipeline runs.
- Logged model metrics and artifacts to MLflow and stored final model weights in a GCS bucket for versioned, production-friendly model management.
- Developed a FastAPI inference service to serve the trained image classifier in real time, exposing a predict endpoint that processes uploaded leaf images and returns disease class probabilities.
- Deployed the serving API on Kubernetes with liveness/readiness probes and autoscaling enabled using HPA for reliable model inference under varying loads.

## TECHNICAL SKILLS

---

- **Languages:** Python,C++,js
- **Operating Systems:** Linux,Windows
- **Cloud Platforms:** AWS,AZURE
- **CI/CD Tools:** Jenkins, Github Actions
- **Container Orchestration:** Docker , Kubernetes
- **IAC:** Terraform